

Results from fitting to reference p-curves

Beth Clarke, Felix Schönbrodt, Angelika Stefan

2026-10-04

While building the simulations (before knowing the actual target p-curve by Clarke et al.), to test our fitting procedure, we used three existing empirical p-curves as targets:

1. **Sotola** ($k = 1,556$ significant p-values): Sotola (2023), Meta-Psychology. [doi](#) → [results](#)
2. **Wetzels** ($k = 593$): Wetzels et al. (2011), Perspectives on Psychological Science. [doi](#) → [results](#)
3. **Simonsohn** ($k = 22$): Simonsohn et al. (2014), JEP: General, Figure 3B. [data](#) → [results](#)

In this document, we report the results of the best-fitting simulations for each of those empirical p-curves including: i) visualisations of the top 5 best fitting simulations to each empirical p-curves for each world; ii) fit indices for the best simulations in each world for each empirical p-curve; iii) the parameter values for the top 10 best fitting p-curves (across all worlds) for each empirical p-curve. :::

Set-up code: Data wrangling and goodness of fit calculations

```
# Combine simulation results from all conditions into one big dataset
devtools::load_all("../..") # load fit p curve with updated plot_pcurves() function
library(rio)
library(tidyverse)
library(ggplot2)
library(kableExtra)

# reference p-curves
pcurves <- import("reference-pcurves.csv")

# load the multiple DV simulation results:
simres_realistic <- import("../sim-results/sim_realistic.csv")
```

```

simres_realistic$condition <- "multDV_realistic"
simres_worst <- import("../sim-results/sim_worst.csv")
simres_worst$condition <- "multDV_worst"
simres_perfect <- import("../sim-results/sim_perfect.csv")
simres_perfect$condition <- "multDV_perfect"
simres_H0 <- import("../sim-results/sim_H0.csv")
simres_H0$condition <- "multDV_H0"

# Delete unnecessary rows for different reporting strategies from simres_perfect
# and simres_H0 (these yield identical results as there is only one p-value)
simres_perfect <- simres_perfect |> filter(strategy == 3)
simres_H0 <- simres_H0 |> filter(strategy == 3)

# load the optional stopping simulation results:
simres_optStop_worst <- import("../sim-results/sim_optStop_worst.csv")
simres_optStop_worst$condition <- "optStop_worst"
simres_optStop_realistic <- import("../sim-results/sim_optStop_realistic.csv")
simres_optStop_realistic$condition <- "optStop_realistic"

simres_multDV <- rbind(simres_realistic, simres_worst, simres_perfect, simres_H0)
simres_multDV$binsum <- rowSums(simres_multDV[, paste0("p", 1:5)])
stopifnot(all(simres_multDV$binsum |> round(6) == 1))
simres_multDV$type <- "multDV"

simres_optStop <- rbind(simres_optStop_realistic, simres_optStop_worst)
simres_optStop$binsum <- rowSums(simres_optStop[, paste0("p", 1:5)])
stopifnot(all(simres_optStop$binsum |> round(6) == 1))
simres_optStop$type <- "optStop"

# Combine simulations from both p-hacking strategies (they have different
# parameter columns)
simres <- bind_rows(simres_multDV, simres_optStop)

# Compute G-squared fit statistics
simres$fit_sotola_g2 <- g2(
  observed = simplify2array(pcurves[1, paste0("p", 1:5)]),
  expected = simres[, paste0("p", 1:5)], n = pcurves$nsig[1], tol=1e-6)

simres$fit_wetzels_g2 <- g2(
  observed = simplify2array(pcurves[2, paste0("p", 1:5)]),
  expected = simres[, paste0("p", 1:5)], n = pcurves$nsig[2], tol=1e-6)

```

```

simres$fit_simonsohn_g2 <- g2(
  observed = simplify2array(pcurves[3, paste0("p", 1:5)]),
  expected = simres[, paste0("p", 1:5)], n = pcurves$nsig[3], tol=1e-6)

# Also compute chi-squared and RMSE for robustness checks
# Chi-squared
simres$fit_sotola_chi2 <- chi2(
  observed = simplify2array(pcurves[1, paste0("p", 1:5)]),
  expected = simres[, paste0("p", 1:5)], n = pcurves$nsig[1], tol=1e-6)

simres$fit_wetzels_chi2 <- chi2(
  observed = simplify2array(pcurves[2, paste0("p", 1:5)]),
  expected = simres[, paste0("p", 1:5)], n = pcurves$nsig[2], tol=1e-6)

simres$fit_simonsohn_chi2 <- chi2(
  observed = simplify2array(pcurves[3, paste0("p", 1:5)]),
  expected = simres[, paste0("p", 1:5)], n = pcurves$nsig[3], tol=1e-6)

# RMSE
simres$fit_sotola_rmse <- rmse(
  observed = simplify2array(pcurves[1, paste0("p", 1:5)]),
  expected = simres[, paste0("p", 1:5)], n = pcurves$nsig[1])

simres$fit_wetzels_rmse <- rmse(
  observed = simplify2array(pcurves[2, paste0("p", 1:5)]),
  expected = simres[, paste0("p", 1:5)], n = pcurves$nsig[2])

simres$fit_simonsohn_rmse <- rmse(
  observed = simplify2array(pcurves[3, paste0("p", 1:5)]),
  expected = simres[, paste0("p", 1:5)], n = pcurves$nsig[3])

# Compute BICs for between-world comparisons (accounting for model complexity)

# Number of free parameters in each simulated "world"
simres$free_parameters <- recode_values(simres$condition,
  "multDV_realistic" ~ 7,
  "multDV_worst" ~ 4,
  "multDV_perfect" ~ 3,
  "multDV_H0" ~ 0,
  "optStop_realistic" ~ 7,
  "optStop_worst" ~ 4)

```

```
# Compute BIC from G-squared and number of free parameters
simres$BIC_sotola <- simres$fit_sotola_g2 + simres$free_parameters * log(pcurves$nsig[1])

simres$BIC_wetzels <- simres$fit_wetzels_g2 + simres$free_parameters * log(pcurves$nsig[1])

simres$BIC_simonsohn <- simres$fit_simonsohn_g2 + simres$free_parameters * log(pcurves$nsig[1])

# Save condition labels for later use
condition_labels <- c(
  multDV_perfect      = "Perfect World",
  multDV_H0           = "Flat World",
  multDV_realistic    = "Realistic World: Multiple DVs",
  multDV_worst        = "Worst World: Multiple DVs",
  optStop_realistic   = "Realistic World: Optional stopping",
  optStop_worst       = "Worst World: Optional stopping"
)
```


1. Target p-curve: Sotola (2023):

Visualisations of best simulations in each world

You can visually inspect the fit of the top 5 (where available) simulations in each world (in blue) to the target p-curve (in black). Note that the Flat World only involved one simulation.


Realistic: Multiples DVs

```
plot_pcurves(simres |> filter(condition == "multDV_realistic"),
  poriginal = simplify2array(pcurves[1, paste0("p", 1:5)]),
  n_best = 5, GOF="g2", n_studies = pcurves$nsig[1],
  tol = 1e-6, title = "Realistic World: Multiple DVs")
```



Realistic: Optional Stopping


```
plot_pcurves(simres |> filter(condition == "optStop_realistic"),
  poriginal = simplify2array(pcurves[1, paste0("p", 1:5)]),
  n_best = 1, GOF="g2", n_studies = pcurves$nsig[1],
  tol = 1e-6, title = "Realistic World: Optional Stopping")
```



Worst: Multiple DVs

```
plot_pcurves(simres |> filter(condition == "multDV_worst"),  
  poriginal = simplify2array(pcurves[1, paste0("p", 1:5)]),  
  n_best = 5, GOF="g2", n_studies = pcurves$nsig[1],  
  tol = 1e-6, title = "Worst Possible World: Multiple DVs")
```


Worst Possible World: Multiple DVs



Worst: Optional Stopping


```
plot_pcurves(simres |> filter(condition == "optStop_worst"),  
  poriginal = simplify2array(pcurves[1, paste0("p", 1:5)]),  
  n_best = 5, GOF="g2", n_studies = pcurves$nsig[1],  
  tol = 1e-6, title = "Worst Possible World: Optional Stopping")
```

Worst Possible World: Optional Stopping




Perfect

```
plot_pcurves(simres |> filter(condition == "multDV_perfect"),  
  poriginal = simplify2array(pcurves[1, paste0("p", 1:5)]),  
  n_best = 5, GOF="g2", n_studies = pcurves$nsig[1],  
  tol = 1e-6, title = "Perfect World")
```



Flat

```
plot_pcurves(simres |> filter(condition == "multDV_H0"),  
  poriginal = simplify2array(pcurves[1, paste0("p", 1:5)]),  
  n_best = 1, GOF="g2", n_studies = pcurves$nsig[1],  
  tol = 1e-6, title = "Flat World")
```



Overview of best fitting simulations in each world

These fit indices reflect the fit of the best-fitting simulation in each condition/ world. Lower numbers indicate better fit. Note that BIC is a transformation of G2 that penalises for model complexity.

```
fit_tab <- simres %>%
  group_by(condition) %>%
  summarise(
    G2    = min(fit_sotola_g2),
    BIC    = min(BIC_sotola),
    Chi2   = min(fit_sotola_chi2),
    RMSE   = min(fit_sotola_rmse)
  )

fit_tab %>%
  mutate(condition = recode(condition, !!!condition_labels)) %>%
  mutate(across(
    -condition,
    ~ cell_spec(sprintf("%.2f", .x), bold = .x == min(.x))
  )) %>%
  kbl(escape = FALSE) %>%
  kable_styling()
```

condition	G2	BIC	Chi2	RMSE
Flat World	139.56	139.56	164.99	0.21
Perfect World	5.62	20.90	5.57	0.02
Realistic World: Multiple DVs	4.55	40.21	4.14	0.02
Worst World: Multiple DVs	20.97	41.34	19.18	0.06
Realistic World: Optional stopping	4.50	40.16	4.17	0.02
Worst World: Optional stopping	160.54	180.91	200.30	0.23

According to G2, the Realistic World: Optional stopping contained the best-fitting simulation to the target p-curve. Once model complexity is accounted for, the Perfect World contained the simulation with the best fit to the target p-curve. Results for chi-squared and RMSE are also provided as robustness checks.

What were the top 10 best fitting simulations (according to G^2)?

These rows reflect the 10 simulations that had the best fit to the target p-curve according to G^2 .

```
simres |> arrange(fit_sotola_g2) |> slice(1:10) |>
  select(condition, nvar, r, d, prop_Hacker, prop_H1, het, strategy, nmin,
         nmax, stepsize, contains("sotola")) |>
  kbl(digits = 2) |>
  kable_styling(latex_options = c("HOLD_position", "scale_down"),
               font_size = 7)
```

condition	nvar	r	d	prop_Hacker	prop_H1	het	strategy	nmin	nmax	stepsize	fit_sotola_g2	fit_sotola_chi2	fit_sotola_rmse	BIC_sotola
optStop_realistic	NA	NA	0.7	0.25	0.2	0.0	NA	444	1063	100	4.50	4.17	0.02	40.16
multDV_realistic	5	0.3	0.6	0.10	0.3	0.2	1	NA	NA	NA	4.55	4.14	0.02	40.21
multDV_realistic	5	0.3	0.6	0.10	0.3	0.2	2	NA	NA	NA	4.55	4.14	0.02	40.21
multDV_realistic	5	0.9	0.8	0.50	0.2	0.2	1	NA	NA	NA	4.82	4.49	0.02	40.47
multDV_realistic	5	0.9	0.8	0.50	0.2	0.2	2	NA	NA	NA	4.82	4.49	0.02	40.47
optStop_realistic	NA	NA	0.7	0.25	0.2	0.2	NA	444	1063	50	4.91	4.42	0.02	40.56
optStop_realistic	NA	NA	0.8	0.10	0.2	0.2	NA	22	1063	100	5.44	4.94	0.02	41.10
multDV_realistic	5	0.3	0.6	0.10	0.3	0.2	3	NA	NA	NA	5.49	4.78	0.03	41.15
multDV_perfect	1	0.0	0.8	0.00	0.2	0.2	3	NA	NA	NA	5.62	5.57	0.02	20.90
optStop_realistic	NA	NA	0.6	0.25	0.2	0.2	NA	444	1063	100	5.71	5.12	0.02	41.37


2. Target p-curve: Wetzels (2011):

Visualisations of best simulations in each world

You can visually inspect the fit of the top 5 (where available) simulations in each world (in blue) to the target p-curve (in black). Note that the Flat World only involved one simulation.

Realistic: Multiple DVs


```
plot_pcurves(simres |> filter(condition == "multDV_realistic"),  
             poriginal = simplify2array(pcurves[2, paste0("p", 1:5)]),  
             n_best = 5, GOF="g2", n_studies = pcurves$nsig[1], tol = 1e-6,  
             title = "Realistic World: Multiple DVs")
```



Realistic: Optional Stopping

```
plot_pcurves(simres |> filter(condition == "optStop_realistic"),  
             poriginal = simplify2array(pcurves[2, paste0("p", 1:5)]),  
             n_best = 1, GOF="g2", n_studies = pcurves$nsig[1], tol = 1e-6,  
             title = "Realistic World: Optional Stopping")
```


Realistic World: Optional Stopping



Worst: Multiple DVs

```
plot_pcurves(simres |> filter(condition == "multDV_worst"),  
  poriginal = simplify2array(pcurves[2, paste0("p", 1:5)]),  
  n_best = 5, GOF="g2",  
  n_studies = pcurves$nsig[1], tol = 1e-6,  
  title = "Worst Possible World: Multiple DVs")
```


Worst Possible World: Multiple DVs



Worst: Optional Stopping


```
plot_pcurves(simres |> filter(condition == "optStop_worst"),  
             poriginal = simplify2array(pcurves[2, paste0("p", 1:5)]),  
             n_best = 5, GOF="g2", n_studies = pcurves$nsig[1],  
             tol = 1e-6, title = "Worst Possible World: Optional Stopping")
```

Worst Possible World: Optional Stopping




Perfect

```
plot_pcurves(simres |> filter(condition == "multDV_perfect"),  
  poriginal = simplify2array(pcurves[2, paste0("p", 1:5)]),  
  n_best = 5, GOF="g2", n_studies = pcurves$nsig[1], tol = 1e-6,  
  title = "Perfect World")
```



Flat

```
plot_pcurves(simres |> filter(condition == "multDV_H0"),  
  poriginal = simplify2array(pcurves[2, paste0("p", 1:5)]),  
  n_best = 1, GOF="g2", n_studies = pcurves$nsig[1], tol = 1e-6,  
  title = "Flat World")
```



Overview of best fitting simulations in each world

These fit indices reflect the fit of the best-fitting simulation in each condition/ world. Lower numbers indicate better fit. Note that BIC is a transformation of G2 that penalises for model complexity.

```
fit_tab <- simres %>%
  group_by(condition) %>%
  summarise(
    G2    = min(fit_wetzels_g2),
    BIC    = min(BIC_wetzels),
    Chi2   = min(fit_wetzels_chi2),
    RMSE   = min(fit_wetzels_rmse)
  )

fit_tab %>%
  mutate(condition = recode(condition, !!!condition_labels)) %>%
  mutate(across(
    -condition,
    ~ cell_spec(sprintf("%.2f", .x), bold = .x == min(.x))
  )) %>%
  kbl(escape = FALSE) %>%
  kable_styling()
```

condition	G2	BIC	Chi2	RMSE
Flat World	555.07	555.07	678.65	0.23
Perfect World	0.57	15.85	0.57	0.00
Realistic World: Multiple DVs	0.13	35.78	0.13	0.00
Worst World: Multiple DVs	26.90	47.28	24.91	0.04
Realistic World: Optional stopping	0.17	35.82	0.16	0.00
Worst World: Optional stopping	634.98	655.36	816.17	0.24

According to G2, the Realistic World: Multiple DVs contained the best-fitting simulation to the target p-curve. Once model complexity is accounted for, the Perfect World contained the simulation with the best fit to the target p-curve. Results for chi-squared and RMSE are also provided as robustness checks.

What were the top 10 best fitting simulations (according to G^2)?

These rows reflect the 10 simulations that had the best fit to the target p-curve according to G^2 .

```
simres |> arrange(fit_wetzels_g2) |> slice(1:10) |>
  select(condition, nvar, r, d, prop_Hacker, prop_H1, het, strategy, nmin,
         nmax, stepsize, contains("wetzels")) |>
  kbl(digits = 2) |>
  kable_styling(latex_options = c("HOLD_position", "scale_down"),
               font_size = 7)
```

condition	nvar	r	d	prop_Hacker	prop_H1	het	strategy	nmin	nmax	stepsize	fit_wetzels_g2	fit_wetzels_chi2	fit_wetzels_rmse	BIC_wetzels
multDV_realistic	2	0.0	0.7	0.50	0.5	0.0	3	NA	NA	NA	0.13	0.13	0	35.78
multDV_realistic	2	0.9	0.6	0.25	0.5	0.2	1	NA	NA	NA	0.16	0.16	0	35.82
multDV_realistic	2	0.9	0.6	0.25	0.5	0.2	2	NA	NA	NA	0.16	0.16	0	35.82
optStop_realistic	NA	NA	0.7	0.10	0.5	0.2	NA	8	444	50	0.17	0.16	0	35.82
multDV_realistic	10	0.0	0.6	0.75	0.5	0.2	1	NA	NA	NA	0.20	0.20	0	35.85
multDV_realistic	10	0.0	0.6	0.75	0.5	0.2	2	NA	NA	NA	0.20	0.20	0	35.85
multDV_realistic	5	0.3	0.6	0.75	0.5	0.0	1	NA	NA	NA	0.20	0.20	0	35.86
multDV_realistic	5	0.3	0.6	0.75	0.5	0.0	2	NA	NA	NA	0.20	0.20	0	35.86
optStop_realistic	NA	NA	0.5	0.50	0.7	0.4	NA	50	444	100	0.22	0.22	0	35.88
multDV_realistic	2	0.9	0.5	0.10	0.7	0.2	1	NA	NA	NA	0.35	0.34	0	36.00


3. Target p-curve: Simonsohn (2014):

Visualisations of best simulations in each world

You can visually inspect the fit of the top 5 (where available) simulations in each world (in blue) to the target p-curve (in black). Note that the Flat World only involved one simulation.

Realistic: Multiple DVs


```
plot_pcurves(simres |> filter(condition == "multDV_realistic"),  
  poriginal = simplify2array(pcurves[3, paste0("p", 1:5)]),  
  n_best = 5, GOF="g2", n_studies = pcurves$nsig[1], tol = 1e-6,  
  title = "Realistic World: Multiple DVs")
```



Realistic: Optional Stopping

```
plot_pcurves(simres |> filter(condition == "optStop_realistic"),  
  poriginal = simplify2array(pcurves[3, paste0("p", 1:5)]),  
  n_best = 1, GOF="g2", n_studies = pcurves$nsig[1], tol = 1e-6,  
  title = "Realistic World: Optional Stopping")
```


Realistic World: Optional Stopping



Worst: Multiple DVs

```
plot_pcurves(simres |> filter(condition == "multDV_worst"),  
             poriginal = simplify2array(pcurves[3, paste0("p", 1:5)]),  
             n_best = 5, GOF="g2", n_studies = pcurves$nsig[1],  
             tol = 1e-6, title = "Worst Possible World: Multiple DVs")
```


Worst Possible World: Multiple DVs



Worst: Optional Stopping


```
plot_pcurves(simres |> filter(condition == "optStop_worst"),  
             poriginal = simplify2array(pcurves[3, paste0("p", 1:5)]),  
             n_best = 5, GOF="g2", n_studies = pcurves$nsig[1], tol = 1e-6,  
             title = "Worst Possible World: Optional Stopping")
```

Worst Possible World: Optional Stopping




Perfect

```
plot_pcurves(simres |> filter(condition == "multDV_perfect"),  
  poriginal = simplify2array(pcurves[3, paste0("p", 1:5)]),  
  n_best = 5, GOF="g2", n_studies = pcurves$nsig[1], tol = 1e-6,  
  title = "Perfect World")
```



Flat

```
plot_pcurves(simres |> filter(condition == "multDV_H0"),  
  poriginal = simplify2array(pcurves[3, paste0("p", 1:5)]),  
  n_best = 1, GOF="g2", n_studies = pcurves$nsig[1],  
  tol = 1e-6, title = "Flat World")
```



Overview of best fitting simulations in each world

These fit indices reflect the fit of the best-fitting simulation in each condition/ world. Lower numbers indicate better fit. Note that BIC is a transformation of G2 that penalises for model complexity.

```
fit_tab <- simres %>%
  group_by(condition) %>%
  summarise(
    G2    = min(fit_simonsohn_g2),
    BIC    = min(BIC_simonsohn),
    Chi2   = min(fit_simonsohn_chi2),
    RMSE   = min(fit_simonsohn_rmse)
  )

fit_tab %>%
  mutate(condition = recode(condition, !!!condition_labels)) %>%
  mutate(across(
    -condition,
    ~ cell_spec(sprintf("%.2f", .x), bold = .x == min(.x))
  )) %>%
  kbl(escape = FALSE) %>%
  kable_styling()
```

condition	G2	BIC	Chi2	RMSE
Flat World	8.05	8.05	7.81	0.12
Perfect World	1.34	16.62	1.21	0.04
Realistic World: Multiple DVs	1.10	36.75	0.97	0.03
Worst World: Multiple DVs	1.72	22.09	1.60	0.04
Realistic World: Optional stopping	1.20	36.85	1.05	0.04
Worst World: Optional stopping	9.26	29.64	9.58	0.13

According to G2, the Realistic World: Multiple DVs contained the best-fitting simulation to the target p-curve. Once model complexity is accounted for, the Flat World contained the simulation with the best fit to the target p-curve. Results for chi-squared and RMSE are also provided as robustness checks.

What were the top 10 best fitting simulations (according to G^2)?

These rows reflect the 10 simulations that had the best fit to the target p-curve according to G^2 .

```
simres |> arrange(fit_simonsohn_g2) |> slice(1:10) |>
  select(condition, nvar, r, d, prop_Hacker, prop_H1, het, strategy, nmin,
    nmax, stepsize, contains("simonsohn")) |>
  kbl(digits = 2) |>
  kable_styling(latex_options = c("HOLD_position", "scale_down"),
    font_size = 7)
```

condition	nvar	r	d	prop_Hacker	prop_H1	het	strategy	nmin	nmax	stepsize	fit_simonsohn_g2	fit_simonsohn_chi2	fit_simonsohn_rmse	BIC_simonsohn
multDV_realistic	70	0.3	0.1	0.25	0.7	0.0	1	NA	NA	NA	1.10	1.13	0.03	36.75
multDV_realistic	70	0.3	0.1	0.25	0.7	0.0	2	NA	NA	NA	1.10	1.13	0.03	36.75
multDV_realistic	70	0.0	0.2	0.10	0.3	0.0	1	NA	NA	NA	1.10	1.10	0.04	36.75
multDV_realistic	70	0.0	0.2	0.10	0.3	0.0	2	NA	NA	NA	1.10	1.10	0.04	36.75
multDV_realistic	150	0.6	0.1	0.25	0.5	0.2	1	NA	NA	NA	1.14	1.01	0.03	36.80
multDV_realistic	150	0.6	0.1	0.25	0.5	0.2	2	NA	NA	NA	1.14	1.01	0.03	36.80
multDV_realistic	2	0.9	0.5	0.25	0.1	0.2	1	NA	NA	NA	1.14	0.97	0.03	36.80
multDV_realistic	2	0.9	0.5	0.25	0.1	0.2	2	NA	NA	NA	1.14	0.97	0.03	36.80
multDV_realistic	70	0.3	0.2	0.10	0.3	0.2	1	NA	NA	NA	1.14	0.98	0.03	36.80
multDV_realistic	70	0.3	0.2	0.10	0.3	0.2	2	NA	NA	NA	1.14	0.98	0.03	36.80